

Gradient Boosting

Recapping Decision Trees

$$T(x, \theta) = \sum_{j=1}^J \gamma_j I(x, R_j)$$

$$\theta = \{R_j, \gamma_j\}_{j=1}^J$$

$$\theta = \arg \max_{\theta} \left(\sum_{j=1}^J \sum_{x_i \in R_j} L(y_i, \gamma_j) \right)$$

Boosted Trees

$$f_m = \sum_{m=1}^M T(x, \theta_m)$$

For Regression Trees

$$\hat{\theta} = \arg \min_{\theta} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + T(x_i, \theta_m))$$

Squared Error Loss : Pick the tree that best predict the residual.

$y_i - f_{m-1}(x_i)$: Target Function

γ_{jm} : average residual error in the R_{jm} region

For **2 - class classification problems** and **Exponential Loss Function**, it becomes exactly same as **AdaBoost solution on decision trees**

Differential Loss Functions

$$L(f) = \sum_{i=1}^N L(y_i, f(x_i))$$

$$\hat{f} = \arg \min_f L(f)$$

$$f = (f(x_1), f(x_2), \dots, f(x_N))$$

$$f_0 = h_0$$

$$f_M = \sum_{m=0}^M h_m \quad h_m \in \mathbb{R}^N$$

– Steepest Descent

$$h_m = -\rho_m g_m$$

$$g_m = \left[\frac{\nabla L(y_i, f(x_i))}{\nabla f(x_i)} \right]_{f(x_i) = f_{m-1}(x_i)}$$

$$\rho_m = \arg \min_{\rho} L(f_{m-1} - \rho g_m)$$

$$f_m = f_{m-1} - \rho_m g_m$$

$g_m \approx$ unconstrained maximal descent direction

$$\hat{\theta}_m = \arg \min_{\theta} \sum_{i=1}^N (-g_{im} - T(x_i; \theta))^2$$

Fit a tree “as close as” possible to gradient direction.

Regression:

$$\frac{1}{2} (y_i - f(x_i))^2 \quad y_i - f(x_i)$$

$$|y_i - f(x_i)| \quad \text{sign}(y_i - f(x_i))$$

Classification:

Deviance $I(y_i = C_k) - P_k(x_i) : (i^{\text{th}} \text{ component})$

$$\hat{\gamma}_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$$

XG BOOST - it's Xtreme..

- XGBoost works as **Newton-Raphson** in function space unlike **gradient boosting** that works as gradient descent in function space.
- A second order **Taylor approximation** is used in the loss function to make the connection to Newton Raphson method.

A generic unregularized XGBoost algorithm is:

Input: training set $\{(x_i, y_i)\}_{i=1}^N$, a differentiable loss function $L(y, F(x))$, a number of weak learners M and a learning rate α .

Algorithm:

1. Initialize model with a constant value:

$$\hat{f}_{(0)}(x) = \arg \min_{\theta} \sum_{i=1}^N L(y_i, \theta).$$

2. For $m = 1$ to M :

1. Compute the 'gradients' and 'hessians':

$$\hat{g}_m(x_i) = \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x)=\hat{f}_{(m-1)}(x)}.$$

$$\hat{h}_m(x_i) = \left[\frac{\partial^2 L(y_i, f(x_i))}{\partial f(x_i)^2} \right]_{f(x)=\hat{f}_{(m-1)}(x)}.$$

XG BOOST continued..

2. Fit a base learner (or weak learner, e.g. tree) using the training set $\left\{ x_i, -\frac{\hat{g}_m(x_i)}{\hat{h}_m(x_i)} \right\}_{i=1}^N$ by solving the optimization problem

below:

$$\hat{\phi}_m = \arg \min_{\phi \in \Phi} \sum_{i=1}^N \frac{1}{2} \hat{h}_m(x_i) \left[\phi(x_i) - \frac{\hat{g}_m(x_i)}{\hat{h}_m(x_i)} \right]^2.$$
$$\hat{f}_m(x) = \alpha \hat{\phi}_m(x).$$

3. Update the model:

$$\hat{f}_{(m)}(x) = \hat{f}_{(m-1)}(x) + \hat{f}_m(x).$$

3. Output $\hat{f}(x) = \hat{f}_{(M)}(x) = \sum_{m=0}^M \hat{f}_m(x).$